

MY AGENT KAFKA

by Model Context Protocol

Guang Zhao

© 2025 NetApp, Inc. All rights reserved.

AGENTIC _____ ?


E.g. Housing agents, car mechanics, doctors

- Interact with high-level language
- Master domain-specific tools
- Hide away domain complexity

KAFKA AGENTIC _____?

Problem: tools are too low-level

TOO LOW-LEVEL



- **✗** "Can you produce X to the topic Y hashed by Z"
- **✓** "How are my sales doing today?"

HIGH-LEVEL PROBLEMS

My Kafka Agentic _____

- Real-time analytics
- Performance tuning

REAL-TIME ANALYTICS



CHEF ASKS...

What are my most popular entrees tonight?

How many do I need to prepare tomorrow?



Write SQL?

Hire a Kafka programmer?


?: Use my Kafka agent?

KAFKA UP FOR ANALYTICS?

"Need to bridge from the operational to the analytical"

-- Confluent engineer at Current '25

DATA MODEL GAP



- Apache Flink
- Kafka Streams
- KSQL (retiring)
- TableFlow (Confluent)
- Iceberg topic (Aiven)
- Kafka table engine (ClickHouse)



: still not easy for chef

CLICKHOUSE = DATA WAREHOUSE + LLM

- Collects 2PB business data
- Analyst:
 - "Which onboarding steps correlate with high 90-day retention"

No longer writing SQLs by hand ... Huge productivity gain

– Mihir Gokhale, PM, Open House '25



: Kafka = Streams + LLM?

PERFORMANCE TUNING



MY KAKFA-ONTAP CONNECTOR

Which ONTAP block size maximizes throughput?

- a. 4KB
- b. 4MB
- c. 4GB



: Guess -> benchmark -> repeat



: Orchestrate for me?

DEMO

- Analyze business data
 - Tune cluster performance
- 👻 Non-deterministic "bugs"?

HOW DOES IT WORK?

HOW COMPLICATED CAN IT BE?


- Input: question
- Compute: ...
 - Programs: models, machine code
 - CPU, storage, network
- Output: answer

How different from "trad coding"?

SYSTEM VIEW

: input

:



: answer

TAKEAWAY

- Agent to solve high-level problems with Kafka
- High-level problems needs more complex tools
- Complexity of tools comes from chosen combination
- LLM is outside of your control*, but tools are within

THOUGHTS?

- <https://github.com/instaclustr/agent-kafka>
- guang@netapp.com
- #opensauce, #open-ai